

Yash-Conqueror /
DSC-Phase3-Project

Code

Issues

Pull requests

Actions

Projects

Security

Insights

[DSC-Phase3-Project](#) / customer_churn_prediction.ipynb 

Yash-source-dev Prediction using SUPPORT VECTOR CLASSIFIER,DECESION TREE CLASSIFIER ...

9357e24 · 2 hours ago



2945 lines (2945 loc) · 297 KB

Preview

Code

Blame

Raw



Introduction

Customer Churn Prediction

Customer attrition or churn, is when customers stop doing business with a company. It can have a significant impact on a company's revenue and it's crucial for businesses to find out the reasons why customers are leaving and take steps to reduce the number of customers leaving. One way to do this is by identifying customer segments that are at risk of leaving, and implementing retention strategies to keep them. Also, by using data and machine learning techniques, companies can predict which customers are likely to leave in the future and take actions to keep them before they decide to leave.

We are going to build a basic model for predicting customer churn using [Telco Customer Churn dataset](#). We are using some classification algorithm to model customers who have left, using Python tools such as pandas for data manipulation and matplotlib for visualizations.

Let's get started.

Steps Involved to Predict Customer Churn

- Importing Libraries
- Loading Dataset
- Exploratory Data Analysis
- Outliers using IQR method
- Cleaning and Transforming Data
 - One-hot Encoding
 - Rearranging Columns
 - Feature Scaling
 - Feature Selection
- Prediction using Logistic Regression
- Prediction using Support Vector Classifier
- Prediction using Decision Tree Classifier
- Prediction using KNN Classifier

Importing Libraries

First of all, we will import known necessary libraries.

In [285...]

```
#import platform
import pandas as pd
import sklearn
```

```
import numpy as np
#import graphviz
import seaborn as sns
import matplotlib
import matplotlib.pyplot as plt
# import plotly.express as px
# import plotly.graph_objects as go

%matplotlib inline
```

Loading Dataset

We use pandas to read the dataset and preprocess it.## Loading Dataset We use pandas to read the dataset and preprocess it.

In [286...]

```
df = pd.read_csv('Data/WA_Fn-UseC_-Telco-Customer-Churn.csv')
df.shape
```

Out[286...]

(7043, 21)

Exploratory Data Analysis

In [287...]

```
df.head()
```

Out[287...]

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService
0	7590-VHVEG	Female	0	Yes	No	1	No
1	5575-GNVDE	Male	0	No	No	34	Yes
2	3668-QPYBK	Male	0	No	No	2	Yes
3	7795-CFOCW	Male	0	No	No	45	No
4	9237-HQITU	Female	0	No	No	2	Yes

5 rows × 21 columns

In [288...]

```
df.tail()
```

Out[288...]

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService
7038	6840-RESVB	Male	0	Yes	Yes	24	
7039	2234-XADUH	Female	0	Yes	Yes	72	

7040	4801-JZAZL	Female	0	Yes	Yes	11
7041	8361-LTMKD	Male	1	Yes	No	4
7042	3186-AJIEK	Male	0	No	No	66

5 rows × 21 columns

In [289...]

```
df.shape
```

Out[289...]

(7043, 21)

We have 2 types of features in the dataset: categorical (two or more values and without any order) and numerical. Most of the feature names are self-explanatory, except for:

- Partner: whether the customer has a partner or not (Yes, No),
- Dependents: whether the customer has dependents or not (Yes, No),
- OnlineBackup: whether the customer has online backup or not (Yes, No, No internet service),
- tenure: number of months the customer has stayed with the company,
- MonthlyCharges: the amount charged to the customer monthly,
- TotalCharges: the total amount charged to the customer.

There are 7043 customers in the dataset and 19 features without customerID (non-informative) and Churn column (target variable). Most of the categorical features have 4 or less unique values.

In [290...]

```
df.size
```

Out[290...]

147903

In [291...]

```
df.dtypes
```

Out[291...]

customerID	object
gender	object
SeniorCitizen	int64
Partner	object
Dependents	object
tenure	int64
PhoneService	object
MultipleLines	object
InternetService	object
OnlineSecurity	object
OnlineBackup	object
DeviceProtection	object
TechSupport	object

```
StreamingTV      object
StreamingMovies   object
Contract          object
PaperlessBilling   object
PaymentMethod     object
MonthlyCharges    float64
TotalCharges      object
Churn             object
dtype: object
```

Totalcharges is given as object datatype but it is float datatype

In [292... df.columns

Out[292... Index(['customerID', 'gender', 'SeniorCitizen', 'Partner', 'Dependents',
 'tenure', 'PhoneService', 'MultipleLines', 'InternetService',
 'OnlineSecurity', 'OnlineBackup', 'DeviceProtection', 'TechSuppor
 t',
 'StreamingTV', 'StreamingMovies', 'Contract', 'PaperlessBilling',
 'PaymentMethod', 'MonthlyCharges', 'TotalCharges', 'Churn'],
 dtype='object')

In [293... df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
 #   Column           Non-Null Count  Dtype  
 ---  -- 
 0   customerID      7043 non-null   object 
 1   gender          7043 non-null   object 
 2   SeniorCitizen   7043 non-null   int64  
 3   Partner         7043 non-null   object 
 4   Dependents     7043 non-null   object 
 5   tenure          7043 non-null   int64  
 6   PhoneService    7043 non-null   object 
 7   MultipleLines   7043 non-null   object 
 8   InternetService 7043 non-null   object 
 9   OnlineSecurity  7043 non-null   object 
 10  OnlineBackup    7043 non-null   object 
 11  DeviceProtection 7043 non-null   object 
 12  TechSupport    7043 non-null   object 
 13  StreamingTV     7043 non-null   object 
 14  StreamingMovies 7043 non-null   object 
 15  Contract        7043 non-null   object 
 16  PaperlessBilling 7043 non-null   object 
 17  PaymentMethod   7043 non-null   object 
 18  MonthlyCharges  7043 non-null   float64
 19  TotalCharges    7043 non-null   object 
 20  Churn           7043 non-null   object 
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB
```

In [294... df.isnull().sum()

Out[294... customerID 0
gender 0
SeniorCitizen 0
Partner 0
Dependents 0
tenure 0
PhoneService 0

```
MultipleLines      0
InternetService   0
OnlineSecurity    0
OnlineBackup       0
DeviceProtection   0
TechSupport        0
StreamingTV        0
StreamingMovies    0
Contract          0
PaperlessBilling   0
PaymentMethod      0
MonthlyCharges     0
TotalCharges       0
Churn              0
dtype: int64
```

In [295... df.duplicated().sum()

Out[295... 0

Basic Data Cleaning:

As we have already observed in above cell that TotalCharges is given as object datatype but it is float datatype. We will fix it here.

In [296... df['TotalCharges'].dtype

Out[296... dtype('O')

In [297... df['TotalCharges']=pd.to_numeric(df['TotalCharges'],errors='coerce')

In [298... df['TotalCharges'].dtype

Out[298... dtype('float64')

In [299... categorical_features = [

```
"gender",
"SeniorCitizen",
"Partner",
"Dependents",
"PhoneService",
"MultipleLines",
"InternetService",
"OnlineSecurity",
"OnlineBackup",
"DeviceProtection",
"TechSupport",
"StreamingTV",
"StreamingMovies",
"Contract",
"PaperlessBilling",
"PaymentMethod",
```

```
]
numerical_features = ["tenure", "MonthlyCharges", "TotalCharges"]
target = "Churn"
```

In [300...]

```
df.skew(numeric_only= True)
```

Out[300...]

SeniorCitizen	1.833633
tenure	0.239540
MonthlyCharges	-0.220524
TotalCharges	0.961642
dtype:	float64

In [301...]

```
numeric_df = df.select_dtypes(include=['number'])
corr_matrix = numeric_df.corr()
corr_matrix
```

Out[301...]

	SeniorCitizen	tenure	MonthlyCharges	TotalCharges
SeniorCitizen	1.000000	0.016567	0.220173	0.102411
tenure	0.016567	1.000000	0.247900	0.825880
MonthlyCharges	0.220173	0.247900	1.000000	0.651065
TotalCharges	0.102411	0.825880	0.651065	1.000000

Feature distribution

We plot distributions for numerical and categorical features to check for outliers and compare feature distributions with target variable.

Numerical features distribution

Numeric summarizing techniques (mean, standard deviation, etc.) don't show us spikes, shapes of distributions and it is hard to observe outliers with it. That is the reason we use histograms.

In [302...]

```
df[numerical_features].describe()
```

Out[302...]

	tenure	MonthlyCharges	TotalCharges
count	7043.000000	7043.000000	7032.000000
mean	32.371149	64.761692	2283.300441
std	24.559481	30.090047	2266.771362
min	0.000000	18.250000	18.800000
25%	9.000000	35.500000	401.450000
50%	29.000000	70.350000	1397.475000
75%	55.000000	89.850000	3794.737500
max	72.000000	118.750000	8684.800000

In [303...]

```
df[numerical_features].hist(bins=30, figsize=(10, 7))
```

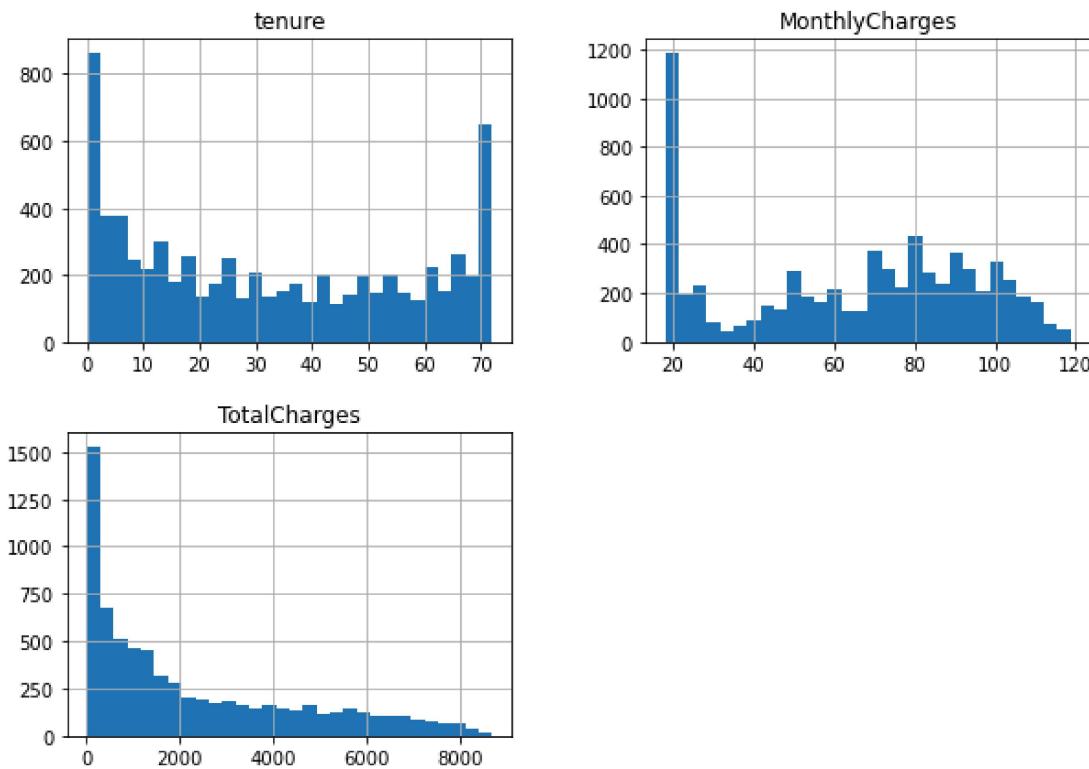
Out[303...]

```
array([[[<AxesSubplot:title={'center': 'tenure'}>],
```

```
DSC-Phase3-Project/customer_churn_prediction.ipynb at main · Yash-Conqueror/DSC-Phase3-Project
<AxesSubplot:title={'center':'MonthlyCharges'}>],  

[<AxesSubplot:title={'center':'TotalCharges'}>, <AxesSubplot:>]],  

dtype=object)
```



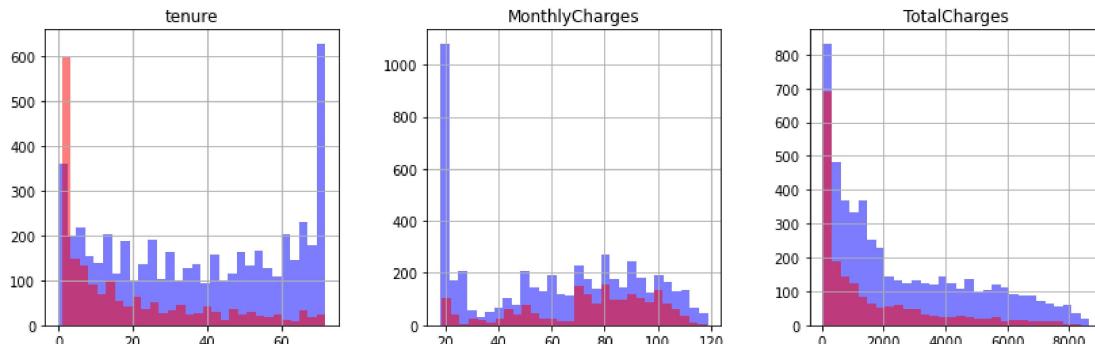
We look at distributions of numerical features in relation to the target variable. We can observe that the greater TotalCharges and tenure are the less is the probability of churn.

In [304...]

```
fig, ax = plt.subplots(1, 3, figsize=(14, 4))
df[df.Churn == "No"][numerical_features].hist(bins=30, color="blue", alpha=0.5)
df[df.Churn == "Yes"][numerical_features].hist(bins=30, color="red", alpha=0.5)
```

Out[304...]

```
array([<AxesSubplot:title={'center':'tenure'}>,
       <AxesSubplot:title={'center':'MonthlyCharges'}>,
       <AxesSubplot:title={'center':'TotalCharges'}>], dtype=object)
```



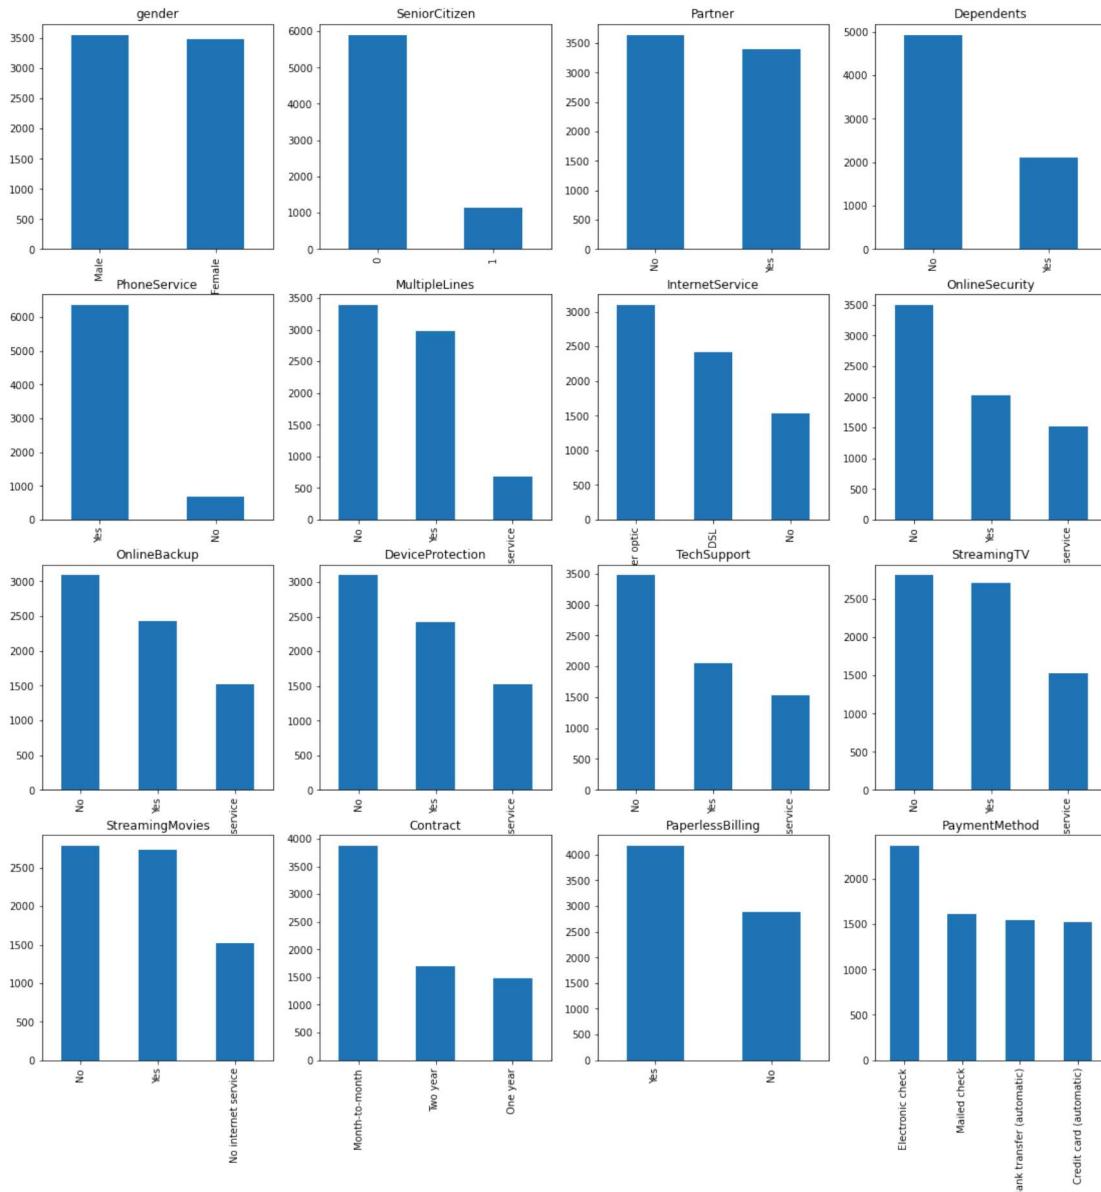
Categorical feature distribution

To analyze categorical features, we use bar charts. We observe that Senior citizens and customers without phone service are less represented in the data.

In [305...]

```
ROWS, COLS = 4, 4
fig, ax = plt.subplots(ROWS,COLS, figsize=(19,19))
row, col = 0, 0,
```

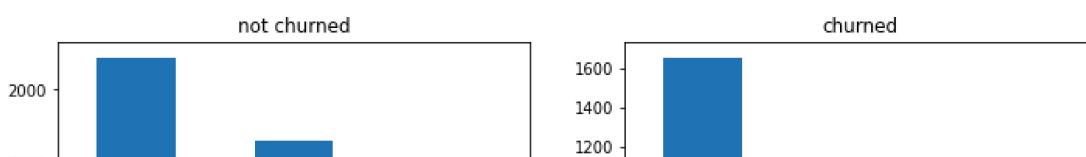
```
for i, categorical_feature in enumerate(categorical_features):
    if col == COLS - 1:
        row += 1
    col = i % COLS
    df[categorical_feature].value_counts().plot(kind='bar', ax=ax[row, col]
```

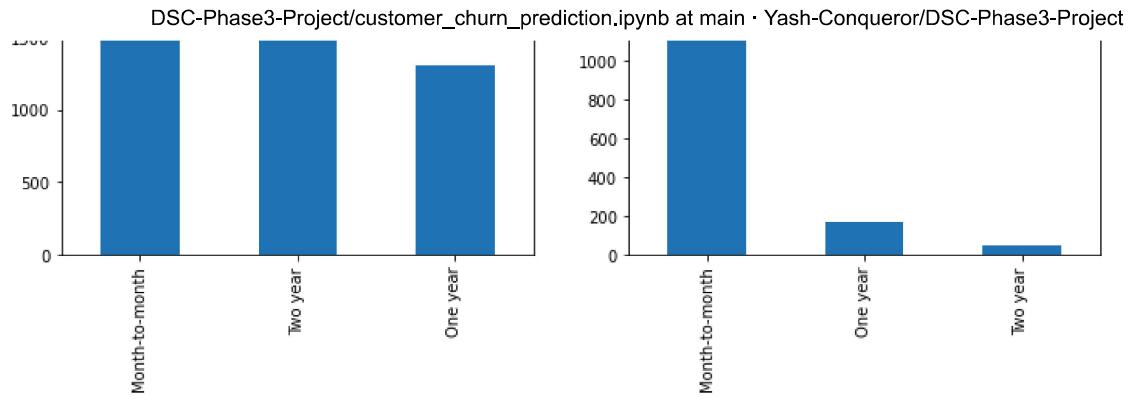


The next step is to look at categorical features in relation to the target variable. We do this only for contract feature. Users who have a month-to-month contract are more likely to churn than users with long term contracts.

```
In [306...]: feature = 'Contract'
fig, ax = plt.subplots(1, 2, figsize=(12, 4))
df[df.Churn == "No"][feature].value_counts().plot(kind='bar', ax=ax[0]).set_ylabel('Count')
df[df.Churn == "Yes"][feature].value_counts().plot(kind='bar', ax=ax[1]).set_ylabel('Count')
```

Out[306...]: Text(0.5, 1.0, 'churned')

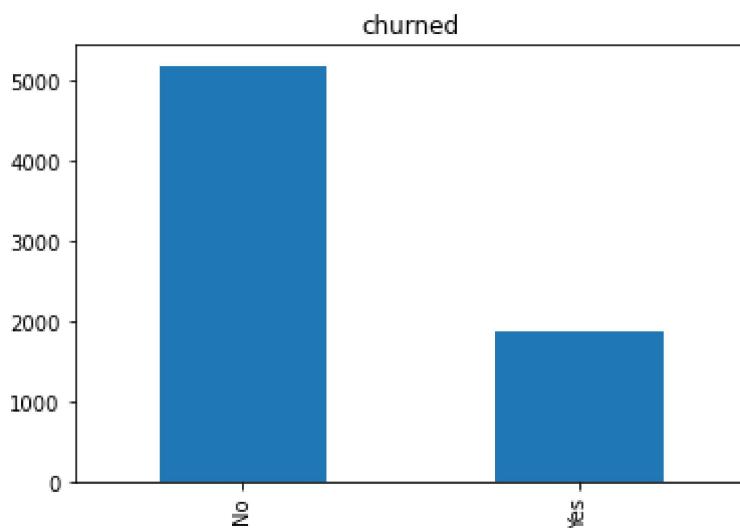




Target variable distribution

```
In [307]: df[target].value_counts().plot(kind='bar').set_title('churned')
```

```
Out[307]: Text(0.5, 1.0, 'churned')
```



Target variable distribution shows that we are dealing with an imbalanced problem as there are many more non-churned as compare to churned users. The model would achieve high accuracy as it would mostly predict majority class - users who didn't churn in our example.

Few things we can do to minimize the influence of imbalanced dataset:

- resample data,
- collect more samples,
- use precision and recall as accuracy metrics.

Outliers Analysis with IQR Method

```
In [308]: x=['tenure', 'MonthlyCharges']
def count_outliers(data,col):
    q1 = data[col].quantile(0.25, interpolation='nearest')
    q2 = data[col].quantile(0.5, interpolation='nearest')
    q3 = data[col].quantile(0.75, interpolation='nearest')
    q4 = data[col].quantile(1, interpolation='nearest')
    IQR = q3 - q1
    global LLP
    global ULP
```

```

LLP = q1 - 1.5*IQR
ULP = q3 + 1.5*IQR
if data[col].min() > LLP and data[col].max() < ULP:
    print("No outliers in",i)
else:
    print("There are outliers in",i)
    x = data[data[col]<LLP][col].size
    y = data[data[col]>ULP][col].size
    a.append(i)
    print('Count of outliers are:',x+y)

global a
a = []
for i in x:
    count_outliers(df,i)

```

No outliers in tenure
No outliers in MonthlyCharges

Cleaning and Transforming Data

In [309... df.drop(['customerID'],axis = 1,inplace = True)

In [310... df.head()

Out[310...

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines
0	Female	0	Yes	No	1	No	No phone service
1	Male	0	No	No	34	Yes	N
2	Male	0	No	No	2	Yes	N
3	Male	0	No	No	45	No	No phone service
4	Female	0	No	No	2	Yes	N

Dropped customerID because it is not needed

On Hot Encoding

In [311... df1=pd.get_dummies(data=df,columns=['gender', 'Partner', 'Dependents', 'PhoneService', 'MultipleLines', 'InternetService', 'OnlineSecurity', 'OnlineBackup', 'DeviceProtection', 'TechSupport', 'StreamingTV', 'StreamingMovies', 'Contract', 'PaperlessBilling', 'PaymentMethod',

In [312... df1.head()

Out[312...]

	SeniorCitizen	tenure	MonthlyCharges	TotalCharges	gender_Male	Partner_Yes
0	0	1	29.85	29.85	0	1
1	0	34	56.95	1889.50	1	0
2	0	2	53.85	108.15	1	0
3	0	45	42.30	1840.75	1	0
4	0	2	70.70	151.65	0	0

5 rows × 31 columns

In [313...]

df1.columns

```
Out[313...]: Index(['SeniorCitizen', 'tenure', 'MonthlyCharges', 'TotalCharges',
       'gender_Male', 'Partner_Yes', 'Dependents_Yes', 'PhoneService_Yes',
       'MultipleLines_No phone service', 'MultipleLines_Yes',
       'InternetService_Fiber optic', 'InternetService_No',
       'OnlineSecurity_No internet service', 'OnlineSecurity_Yes',
       'OnlineBackup_No internet service', 'OnlineBackup_Yes',
       'DeviceProtection_No internet service', 'DeviceProtection_Yes',
       'TechSupport_No internet service', 'TechSupport_Yes',
       'StreamingTV_No internet service', 'StreamingTV_Yes',
       'StreamingMovies_No internet service', 'StreamingMovies_Yes',
       'Contract_One year', 'Contract_Two year', 'PaperlessBilling_Yes',
       'PaymentMethod_Credit card (automatic)',
       'PaymentMethod_Electronic check', 'PaymentMethod_Mailed check',
       'Churn_Yes'],
      dtype='object')
```

Rearranging Columns

In [314...]

```
df1 = df1[['SeniorCitizen', 'tenure', 'MonthlyCharges', 'TotalCharges',
       'gender_Male', 'Partner_Yes', 'Dependents_Yes',
       'PhoneService_Yes', 'MultipleLines_No phone service',
       'MultipleLines_Yes', 'InternetService_Fiber optic',
       'InternetService_No', 'OnlineSecurity_No internet service',
       'OnlineSecurity_Yes', 'OnlineBackup_No internet service',
       'OnlineBackup_Yes', 'DeviceProtection_No internet service',
       'DeviceProtection_Yes', 'TechSupport_No internet service',
       'TechSupport_Yes', 'StreamingTV_No internet service', 'StreamingTV_Yes',
       'StreamingMovies_No internet service', 'StreamingMovies_Yes',
       'Contract_One year', 'Contract_Two year', 'PaperlessBilling_Yes',
       'PaymentMethod_Credit card (automatic)',
       'PaymentMethod_Electronic check', 'PaymentMethod_Mailed check', 'Chu
```

In [315...]

df1.head()

Out[315...]

	SeniorCitizen	tenure	MonthlyCharges	TotalCharges	gender_Male	Partner_Yes
0	0	1	29.85	29.85	0	1

1	0	34	56.95	1889.50	1	0
2	0	2	53.85	108.15	1	0
3	0	45	42.30	1840.75	1	0
4	0	2	70.70	151.65	0	0

5 rows × 31 columns

◀ ▶

In [316]: df1.shape

Out[316]: (7043, 31)

```
from sklearn.impute import SimpleImputer

# The imputer will replace missing values with the mean of the non-missing

imputer = SimpleImputer(missing_values=np.nan, strategy="mean")

df1.TotalCharges = imputer.fit_transform(df1["TotalCharges"].values.reshape(-1, 1))
```

◀ ▶

Feature Scaling

In [318]:

```
from sklearn.preprocessing import StandardScaler
scaler=StandardScaler()
```

In [319]:

```
scaler.fit(df1.drop(['Churn_Yes'],axis = 1))
scaled_features = scaler.transform(df1.drop('Churn_Yes',axis = 1))
```

Feature Selection

In [320]:

```
from sklearn.model_selection import train_test_split
X = scaled_features
y= df1['Churn_Yes']
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size = 0.3,random_state=42)
```

◀ ▶

Prediction using Logistic Regression

In [321]:

```
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report,accuracy_score ,confusion_matrix

logmodel = LogisticRegression()
logmodel.fit(X_train,y_train)
```

◀ ▶

Out[321]: LogisticRegression()

```
In [322... predLR=logmodel.predict(X_test)
```

```
In [323... predLR
```

```
Out[323... array([0, 0, 0, ..., 0, 0, 0], dtype=uint8)
```

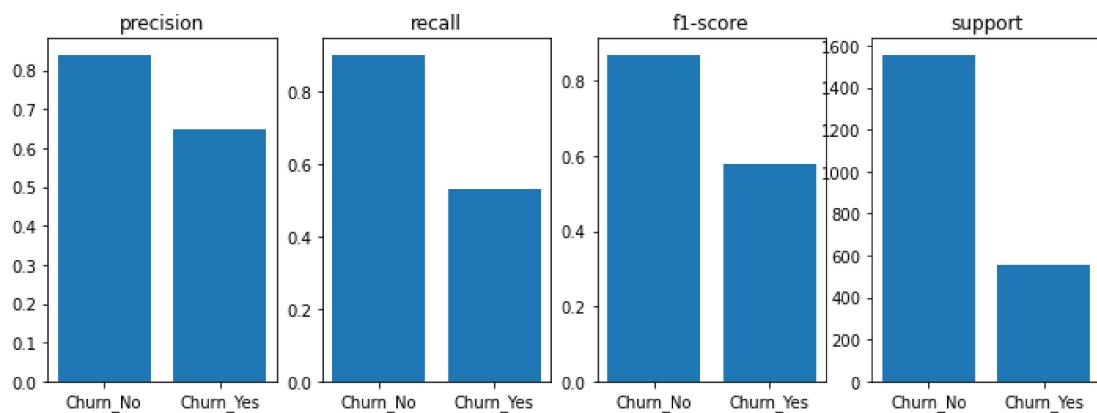
```
In [324... y_test
```

```
Out[324... 5616    0  
2937    0  
1355    0  
5441    1  
3333    0  
..  
2797    1  
412     0  
174     0  
5761    0  
5895    0  
Name: Churn_Yes, Length: 2113, dtype: uint8
```

```
In [325... print(classification_report(y_test,predLR))
```

	precision	recall	f1-score	support
0	0.84	0.90	0.87	1557
1	0.65	0.53	0.58	556
accuracy			0.80	2113
macro avg	0.74	0.71	0.73	2113
weighted avg	0.79	0.80	0.79	2113

```
In [326... # calculate the classification report  
report = classification_report(y_test, predLR, target_names=['Churn_No', 'Churn_Yes'])  
  
# split the report into lines  
lines = report.split('\n')  
  
# split each line into parts  
parts = [line.split() for line in lines[2:-5]]  
  
# extract the metrics for each class  
class_metrics = dict()  
for part in parts:  
    class_metrics[part[0]] = {'precision': float(part[1]), 'recall': float(part[2]), 'f1-score': float(part[3]), 'support': int(part[4])}  
  
# create a bar chart for each metric  
fig, ax = plt.subplots(1, 4, figsize=(12, 4))  
metrics = ['precision', 'recall', 'f1-score', 'support']  
for i, metric in enumerate(metrics):  
    ax[i].bar(class_metrics.keys(), [class_metrics[key][metric] for key in class_metrics])  
    ax[i].set_title(metric)  
  
# display the plot  
plt.show()
```



```
In [327...]: confusion_matrix_LR = confusion_matrix(y_test, predLR)
```

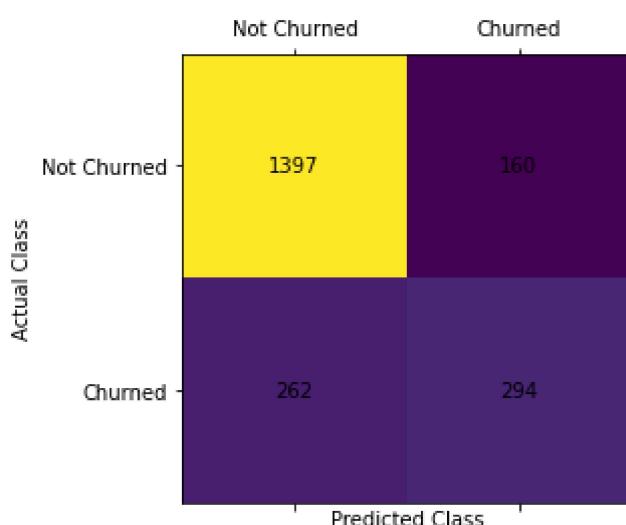
```
In [328...]: # create a heatmap of the matrix using matshow()

plt.matshow(confusion_matrix(y_test, predLR))

# add Labels for the x and y axes
plt.xlabel('Predicted Class')
plt.ylabel('Actual Class')

for i in range(2):
    for j in range(2):
        plt.text(j, i, confusion_matrix_LR[i, j], ha='center', va='center')

# Add custom Labels for x and y ticks
plt.xticks([0, 1], ["Not Churned", "Churned"])
plt.yticks([0, 1], ["Not Churned", "Churned"])
plt.show()
```



```
In [329...]: logmodel.score(X_train, y_train)
```

```
Out[329...]: 0.8062880324543611
```

```
In [330...]: accuracy_score(y_test, predLR)
```

```
Out[330...]: 0.8002839564600095
```

Prediction using Support Vector Classifier

```
In [331...]: from sklearn.svm import SVC
```

```
svc = SVC()
svc.fit(X_train, y_train)
y_pred_svc = svc.predict(X_test)
```

```
In [332...]: print(classification_report(y_test, y_pred_svc))
```

	precision	recall	f1-score	support
0	0.83	0.92	0.87	1557
1	0.67	0.48	0.56	556
accuracy			0.80	2113
macro avg	0.75	0.70	0.71	2113
weighted avg	0.79	0.80	0.79	2113

```
In [333...]: confusion_matrix_svc = confusion_matrix(y_test, y_pred_svc)
```

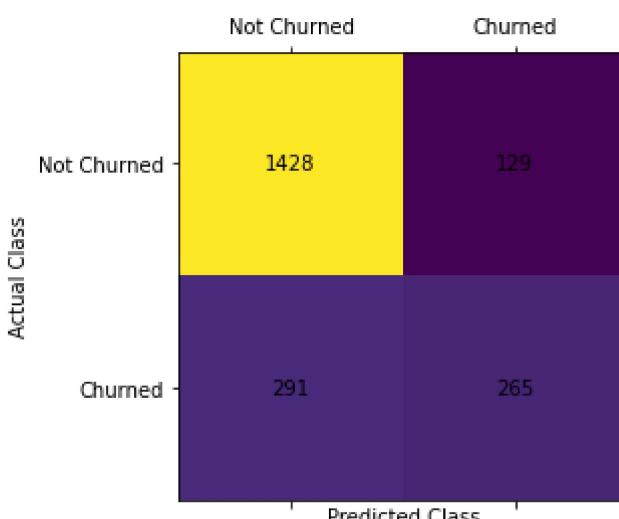
```
In [334...]: # create a heatmap of the matrix using matshow()
```

```
plt.matshow(confusion_matrix_svc)

# add labels for the x and y axes
plt.xlabel('Predicted Class')
plt.ylabel('Actual Class')

for i in range(2):
    for j in range(2):
        plt.text(j, i, confusion_matrix_svc[i, j], ha='center', va='center')

# Add custom labels for x and y ticks
plt.xticks([0, 1], ["Not Churned", "Churned"])
plt.yticks([0, 1], ["Not Churned", "Churned"])
plt.show()
```



```
In [335...]: svc.score(X_train,y_train)
```

```
Out[335...]: 0.8170385395537525
```

Prediction using Decision Tree Classifier

```
In [336...]: from sklearn.tree import DecisionTreeClassifier
dtc = DecisionTreeClassifier()
dtc.fit(X_train, y_train)
y_pred_dtc = dtc.predict(X_test)
```

```
In [337...]: print(classification_report(y_test, y_pred_dtc))
```

	precision	recall	f1-score	support
0	0.81	0.81	0.81	1557
1	0.47	0.47	0.47	556
accuracy			0.72	2113
macro avg	0.64	0.64	0.64	2113
weighted avg	0.72	0.72	0.72	2113

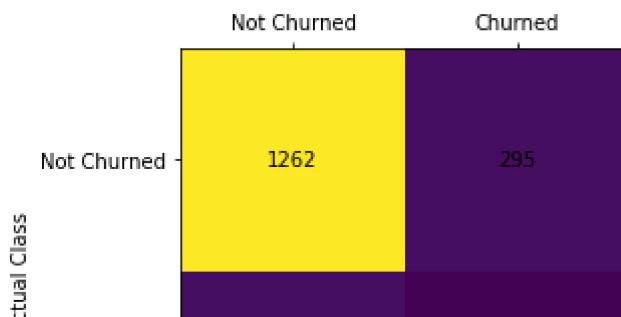
```
In [338...]: confusion_matrix_dtc = confusion_matrix(y_test, y_pred_dtc)
```

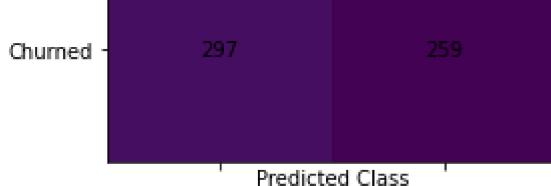
```
# create a heatmap of the matrix using matshow()
plt.matshow(confusion_matrix_dtc)

# add Labels for the x and y axes
plt.xlabel('Predicted Class')
plt.ylabel('Actual Class')

for i in range(2):
    for j in range(2):
        plt.text(j, i, confusion_matrix_dtc[i, j], ha='center', va='center')

# Add custom labels for x and y ticks
plt.xticks([0, 1], ["Not Churned", "Churned"])
plt.yticks([0, 1], ["Not Churned", "Churned"])
plt.show()
```





In [340...]: `dtc.score(X_train,y_train)`

Out[340...]: 0.9987829614604462

In [341...]: `accuracy_score(y_test, y_pred_dtc)`

Out[341...]: 0.7198296261239944

Prediction using KNN Classifier

In [342...]: `from sklearn.neighbors import KNeighborsClassifier`

```
knn = KNeighborsClassifier(n_neighbors = 30)
knn.fit(X_train,y_train)
```

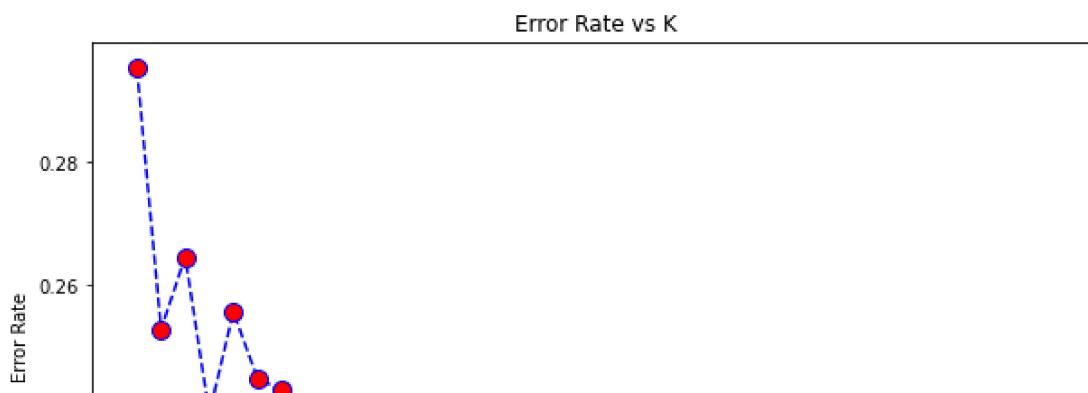
Out[342...]: `KNeighborsClassifier(n_neighbors=30)`

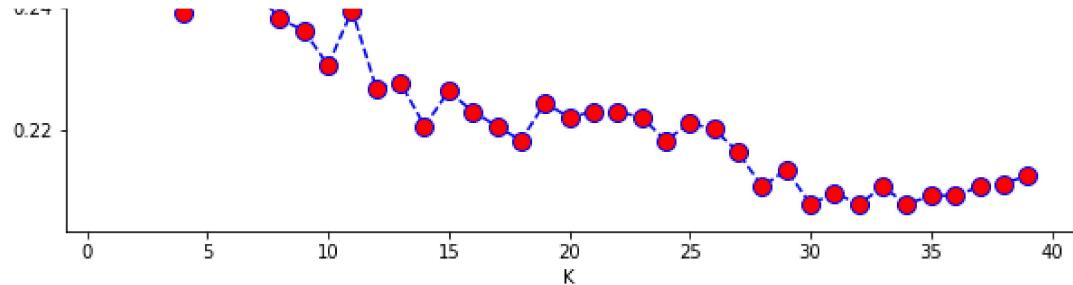
In [343...]: `pred_knn = knn.predict(X_test)`

```
error_rate= []
for i in range(1,40):
    knn = KNeighborsClassifier(n_neighbors = i)
    knn.fit(X_train,y_train)
    pred_i = knn.predict(X_test)
    error_rate.append(np.mean(pred_i != y_test))
```

In [345...]: `plt.figure(figsize = (10,6))
plt.plot(range(1,40),error_rate,color = 'blue',linestyle = '--',marker = 'o')
plt.title('Error Rate vs K')
plt.xlabel('K')
plt.ylabel('Error Rate')`

Out[345...]: `Text(0, 0.5, 'Error Rate')`





In [346...]

```
print(classification_report(y_test,pred_knn))
```

	precision	recall	f1-score	support
0	0.84	0.88	0.86	1557
1	0.62	0.55	0.58	556
accuracy			0.79	2113
macro avg	0.73	0.71	0.72	2113
weighted avg	0.79	0.79	0.79	2113

In [347...]

```
confusion_matrix_knn = confusion_matrix(y_test,pred_knn)
```

In [348...]

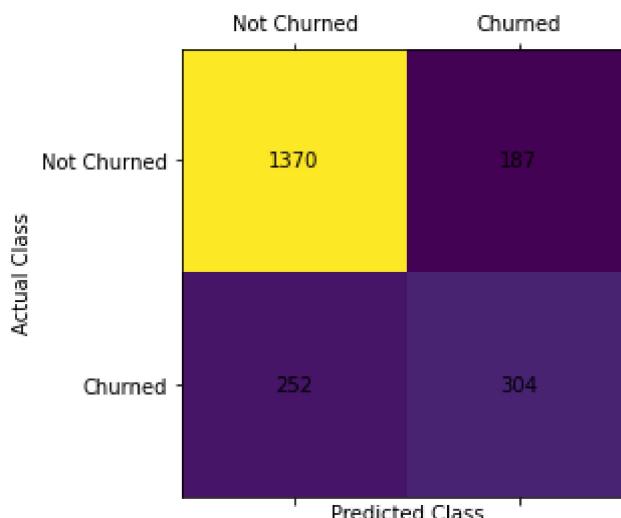
```
# create a heatmap of the matrix using matshow()

plt.matshow(confusion_matrix_knn)

# add labels for the x and y axes
plt.xlabel('Predicted Class')
plt.ylabel('Actual Class')

for i in range(2):
    for j in range(2):
        plt.text(j, i, confusion_matrix_knn[i, j], ha='center', va='center')

# Add custom labels for x and y ticks
plt.xticks([0, 1], ["Not Churned", "Churned"])
plt.yticks([0, 1], ["Not Churned", "Churned"])
plt.show()
```



In [349...]

```
knn.score(X_train,y_train)
```

```
Out[349...]: 0.8008113590263691
```

```
In [350...]: accuracy_score(y_test, pred_knn)
```